

Helix — Self-Verifying Agent

Every coding agent's weakness is writing code that looks right but doesn't run. Helix can actually run it. After a build, the agent runs the project in a real sandbox, reads the errors, fixes them, and re-runs — until it works or it honestly says what's broken. This completes Plan → Build → Verify.

1 · PLAN

Propose

Read-only — the agent lays out a numbered plan you approve.



2 · BUILD

Write the code

The agent writes the files into your workspace.



3 · VERIFY

Run & prove it

Runs the build in the sandbox, reads errors, fixes, re-runs.

↑ on failure, the fix loops back into Build automatically (capped) ↑

What you see in chat

- ✓ Built 4 files
- ⚙ Verifying — `npm run build...`
- ✗ `TypeError in app.tsx:12 — fixing...`
- 🔄 Re-running `npm run build...`
- Verified · build is clean

✓ Verified (auto-fixed) — click to view the run log

How it works

Real execution. Verify runs the project in Helix's sandbox VM via the same engine the live preview uses — not a guess.

Smart command. Picks `npm run build` or `npm test` from your `package.json`; plain HTML and script-less projects skip for free.

Auto-fix loop. On a non-zero exit, the build error is fed back to the agent to fix, then re-run — bounded so it never spins.

Honest by default. If it can't be fixed, you get **Couldn't verify** with the exact log — never a false "done."

Two ways in. An auto-verify toggle next to Build/Plan, or a per-message **Verify build** button to check on demand.